



Trust-based autonomy to prevent agentic operational loss.

CONTROL TOWER is a mobile-first B2B SaaS control system for autonomous AI agents.

\$500K PRE-SEED

alpha-ready → beta/release

AI agents are gaining action, but not accountability.

Problem

**The next AI risk is not bad answers.
It is bad actions.**

Teams need control before agents create operational loss.

NEW FAILURE MODE

sending

updating

routing

approving

accessing

spending

Wrong output becomes operational consequence when agents act across real systems.

From chatbots to autonomous agents.

WHY NOW



More autonomy means more need for trust-based control.

The risk is no longer theoretical.

WHY NOW

The risk is no longer theoretical.

RISK 01

Uncontrolled action

Agent acts without
safeguards or approval
Production systems at risk

RISK 02

Uncontrolled access

Agent reaches systems,
data, and credentials
No audit trail, no boundaries

RISK 03


Uncontrolled spend

Hundreds of millions burned
in a single month
No caps, no controls, no alerts

Autonomous systems create loss when action, access, and spend move faster than control.

CONTROL TOWER brings trust-based autonomy to AI agents.


Solution



Agents can act. Operators stay in control.

How CONTROL TOWER works


Product



A reviewable control flow before autonomous action reaches the business.

ARCHE has two product directions.

Platform



Start where agent risk is already operational.

Go-to-market

CONTROL TOWER is built for businesses using AI agents.

The first wedge is companies already deploying agents into workflows where mistakes can trigger operational loss: sending, updating, routing, approving, accessing, or spending across business systems.

Any business using agents is a potential customer. CONTROL TOWER starts where agent risk is already operational.

INITIAL TARGETS

- AI automation agencies
- Internal ops teams
- RevOps / sales automation
- Customer support automation
- Operations teams
- Enterprise AI teams
- Agent builders managing deployments

A large market forming around agent control.

Market

TAM **\$650B+**

AI software, agentic systems, governance & trust infrastructure

SAM **\$100B+**

Agent oversight, AI governance, AI TRiSM & workflow control

SOM **\$100M**

Initial wedge: teams already deploying AI agents

Source: ArchePersona Market Research & Platform Expansion Analysis, June 2026.

B2B SaaS for trust-based autonomy.

Business model



Simple SaaS economics with clear expansion.

Financials / unit economics

CONTROL TOWER starts at \$499/month and scales to \$2,000/month for teams managing up to 250 agents.

BASE TIER

\$499/mo

\$5,988 ARR

Up to 5 agents
Approval + audit layer

SCALED TIER

\$2,000/mo

\$24,000 ARR

Up to 250 agents
Full governance stack

ARR PATH AT SCALED TIER

10 customers	\$240K ARR
25 customers	\$600K ARR
50 customers	\$1.2M ARR
100 customers	\$2.4M ARR

PRE-SEED FINANCIAL MILESTONE

3-5 design partners | 1-2 paid pilots or LOIs | beta/release-ready product | seed-ready revenue narrative

The \$500K raise is designed to convert alpha-ready product into paid pilot signal and seed-ready financial proof.

The gap is between agent intent and business action.

Competition / differentiation

Existing tools solve pieces of the agent stack. CONTROL TOWER focuses on the operational checkpoint before autonomous action.

Category	Typical focus	Where CONTROL TOWER differs
Agent builders	Create and deploy agents	Controls what agents can do before action
Observability tools	Monitor traces and behavior	Adds approval, blocking, escalation, and audit workflow
AI governance	Policy, compliance, model risk	Focuses on runtime agent action and operator control
Security / identity	Access, auth, permissions	Adds human-in-the-loop action approval and consequence records
Internal workflows	Manual review processes	Purpose-built for agentic operational risk

CONTROL TOWER does not just monitor agents. It gives teams a trust checkpoint before autonomous actions reach the business.

ARCHE is not pre-product.

Traction / proof

BRUNEL is live and demonstrates ARCHE's behavior engine in real interaction.

CONTROL TOWER is alpha-ready as the first commercial product for trust-based autonomy.

Proof is currently product signal, demand signal, and market signal — not revenue claims.

PROOF STACK

- BRUNEL live demo available
- Founder-built behavior engine
- CONTROL TOWER alpha-ready
- 20K requests in 48 hours
- Visible behavioral difference
- Early investor interest after demo

Founder-built behavior engine, now commercializing trust-based autonomy.

Founder / team



Darren Hall

Founder, CEO & Chief Architect

- Founder-built behavior engine
- BRUNEL live demo available
- CONTROL TOWER alpha-ready
- 30 years in distributed systems & IT
- Military background — radiation systems technology

ADVISORS

- Strategy & product
- Technical execution
- Peer review & academia
- Seasoned business ownership

POST-RAISE HIRES

- CTO / technical lead
- Junior developer

From alpha-ready to beta/release.

Roadmap

NOW

CONTROL TOWER
alpha-ready

BRUNEL live
behavior proof

Mobile-first
trust workflow

30 DAYS

Demo to investors
+ pilot targets

3–5 design partner
conversations

IP / patent review

Legal formation

CTO search begins

60 DAYS

Alpha → beta
candidate

Design partner
feedback loop

Strengthen approval,
audit + escalation

Junior dev hiring

90 DAYS

Beta / pilot release

1–2 design partners
→ paid pilots / LOIs

Usage + customer
proof produced

Seed-ready investor
package prepared

Raising \$500K pre-seed.

Use of funds

Raising \$500K pre-seed to move CONTROL TOWER from alpha-ready into beta/release, secure design partners, protect core IP, and build the foundation for scale.

USE OF FUNDS

- CONTROL TOWER beta/release buildout
- Patent/IP protection
- Incorporation and legal formation
- CTO / technical lead hire
- Junior developer hire
- Design partner and pilot development
- Infrastructure and deployment stability
- Founder runway

DISCIPLINED RAISE

Enough to convert alpha-ready product into beta/release — not an oversized round before traction is converted.

Target outcomes: design partners, paid pilots or LOIs, protected IP, first technical hires, and seed-ready proof.



Trust-based autonomy to prevent agentic operational loss.

Raising \$500K pre-seed to move CONTROL TOWER from alpha-ready into beta/release, secure design partners, protect core IP, and build the foundation for scale.

\$500K PRE-SEED

Backup: 15-market platform expansion map.

Backup

We enter through CONTROL TOWER. Here is how ARCHE expands.

CONTROL TOWER MARKETS

- AI agents
- Agentic AI
- Workflow automation
- AI governance
- AI TRiSM
- Compliance automation
- Cybersecurity / AI safety
- Enterprise operations
- Customer support automation
- RevOps / sales automation

CHIMERA MARKETS

- AI companions / persistent assistants
- Character-based AI agents
- Branded AI agents / corporate personas
- Digital identity / persona systems

Source: ArchePersona Market Research & Platform Expansion Analysis, June 2026.

Backup: ARCHEngine modules.

Technical diligence

ARCHEngine

V-HOLD

TRUST BALANCER

DELIVERY ROOM

MDSA

ARKHIVE

INVESTOR-SAFE TRANSLATION

V-HOLD → action control

**TRUST
BALANCER** → risk and confidence regulation

DELIVERY ROOM → controlled response/output layer

MDSA → behavioral signal analysis

ARKHIVE → memory/archive layer

No deep architecture, system prompts, repo structure, or internal mechanisms in the main deck.

Backup: Market risk evidence.

Sources

The market research package includes risk statistics around enterprise AI losses, agentic AI cancellations, failed production deployments, and enterprise AI spend producing limited measurable results.

EVIDENCE CATEGORIES

- AI agent project failures
- Cost overruns
- Production system incidents
- Governance gaps
- Uncontrolled usage / spend
- Enterprise AI spend with weak ROI

Source: ArchePersona Market Research & Platform Expansion Analysis, June 2026.

Backup: competitive landscape examples.

Competitive diligence

Representative categories only. CONTROL TOWER is positioned at the operational checkpoint before agent action.

Category	Examples	Position
Agent builders	LangChain, CrewAI, AutoGen, OpenAI Agents / Assistants	Build and orchestrate agents
Observability	LangSmith, Langfuse, Arize Phoenix, AgentOps-style tools	Monitor traces, events, and behavior
Governance	Credo AI, Holistic AI, IBM watsonx.governance	Policy, compliance, and model risk
Security / authorization	Arcade.dev, identity/security platforms, MCP/A2A controls	Access, permissions, runtime policy
CONTROL TOWER	Mobile-first trust-based autonomy	Approve, block, escalate, intervene, and audit before action

Strategic advantage: timing + control surface. CONTROL TOWER operates where intent becomes consequence.